

Homework 8 solution

Theory question

1. Zhang's algorithm for camera calibration involves identifying the intrinsic camera parameters from images of a known planar pattern (usually a checkerboard). The absolute conic Ω_∞ represents sets of points at infinity in the domain space, that it is invariant to Euclidian transformation and useful for obtaining intrinsic camera properties. When the plane π that contains the calibration pattern) intersects Ω_∞ , it does so at these two circular points. These circular points provide constraints that help estimate the camera's intrinsic parameters without requiring knowledge of the camera's position or orientation.
2. In homogenous coordinate, the absolute conic is represented as

$$\Omega_\infty = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

When a camera observes Ω_∞ , it is projected as the image of the absolute conic ω . The transformation from the domain coordinate to the image coordinate is governed by the intrinsic camera matrix K :

$$K = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

Where f_x and f_y are the focal lengths along the x and y axes, s is the skew coefficient and (c_x, c_y) is the principal point of the camera.

Projecting Ω_∞ through K gives the image of the conic ω on the image plane, resulting in,

$$\omega = (KK^T)^{-1}$$

(ii) ω is an imaginary conic in the image plane, meaning it does not correspond to any real points or pixel locations in the image. The Absolute Conic Ω_∞ in 3D space lies entirely in the "plane at infinity." When this conic is projected onto the image plane, it remains an imaginary entity because it represents directions rather than actual locations. Mathematically, the ω has complex conjugate points as its "points" of intersection in the projective plane, meaning it has no real solutions in the Euclidean sense.

Specifically, the matrix ω defines a quadratic form:

$$x^T \omega x = 0$$

Where x is a point in homogenous coordinates $(x, y, 1)^T$. For real values of x and y , this equation does not yield any real solutions because ω is an imaginary conic. Therefore, no real points (x, y) on the image plane satisfy this equation.

Canny edge detection, Hough line and intersection labelling

Canny edge detection was used to detect the edges in the checkerboard patterns before Hough Line Transformation, was used to identify and averages vertical and horizontal lines. The averaged lines were refined and used to determine intersections for corner extraction. The resulting visualization from this process can be seen in Figures 1, 3, 8 and 10. The homographies for all image views which map an image's points to the world coordinate system was also calculated.

Intrinsic and extrinsic camera parameters

Using the homography matrices derived from the previous section, the intrinsic parameters of a camera was obtained. First, a set of linear constraints from each homography was used to solve the image of the absolute conic. This matrix was used to extract intrinsic parameters such as focal lengths, skew, and principal point. The extrinsic camera parameters (rotation matrices and translation vectors), were derived using homographies and the intrinsic parameters. The rotation matrix was conditioned to make it orthonormal. These parameters are presented in later sections of this report. Using the intrinsic properties, the images were reprojected. The results of initial reprojection can be seen in Figures 2, 4, 9 and 11.

To get intrinsic parameters using Zhang's method, the relationship between a 3D point (X, Y, Z) in the world coordinates and its projection onto the image plane $x = (x, y, w)^T$ can be represented using a homography when $Z = 0$

$$\begin{matrix} x & X \\ \lambda y & = H Y \\ 1 & 1 \end{matrix}$$

H is the homography that maps points in image to the domain space. λ is a scaling factor

$$H = K[r_1, r_2, t]$$

K is the camera matrix, r_1 and r_2 are the rotational matrix and t is the translation vector

$$K = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

From multiple views of the calibration pattern, we obtain several homography matrices H_i . Using these homographies, we can estimate the intrinsic parameters by setting up the following constraints such that we have symmetric matrix B :

$$B = K^{-T} K^{-1} = \begin{bmatrix} B_{11} & B_{12} & B_{13} \\ B_{12} & B_{22} & B_{23} \\ B_{13} & B_{23} & B_{33} \end{bmatrix}$$

For each homography,

$$B = h_{i1}^T B h_{i2} = 0$$

$$h_{i1}^T B h_{i1} = h_{i2}^T B h_{i2}$$

These constraints can be arranged as a system of linear equations in terms of the elements of \mathbf{B} , which are related to the intrinsic parameters.

After solving for B, we can now get the elements of the K matrix using the relationships below

$$c_x = \frac{B_{13}}{B_{11}}$$

$$c_y = \frac{B_{23}}{B_{22}}$$

$$f_x = \sqrt{\frac{B_{11}}{B_{33} - c_x^2}}$$

$$s = -\frac{B_{12}}{B_{11}}$$

Given a homography matrix $H = [h_1 \ h_2 \ h_3]$ and knowing the intrinsic matrix K, we aim to compute the rotation matrix R and the translation vector t

First, a scale factor because the left side is homogeneous, and we must account for a scale factor ξ to make both sides equivalent

$$\xi = \frac{1}{||K^{-1}h_1||}$$

We compute the rotation vectors r_1, r_2 and r_3

$$r_1 = \xi K^{-1}h_1$$

$$r_2 = \xi K^{-1}h_2$$

$$r_3 = r_1 \times r_2$$

$$t = \xi K^{-1}h_3$$

Levenberg-Marquardt optimization

It can be seen that the initial reprojection has a lot of errors visually. To refine the calibration parameters, Levenberg-Marquardt non-linear optimization was used for better alignment between model and observed points. Figures 5, 6, 12 and 13 show the visual results from this optimization. The new sets of extrinsic properties and camera matrix are also presented in the report.

Radial distortion correction

This radial distortion in projected image points was corrected by adjusting them based on their radial distance from the distortion center. This correction results in undistorted points, aligning the image projection with real-world geometry more accurately. Figures 5, 6, 12 and 13 show the visual results after radial distortion removal and the radial distortion coefficients were presented in the report.

Reprojection error mean and variance

Reprojection analysis was performed to evaluate the accuracy of a camera model by projecting 3D world coordinates onto 2D image coordinates and measuring the error. The mean and variance of reprojection errors by comparing actual image points with the projected points are presented in Tables 1 and 2.

Camera pose

The 3D representation of all views from the cameras are presented in Figures 7 and 14.

Custom dataset output

Image 4 output

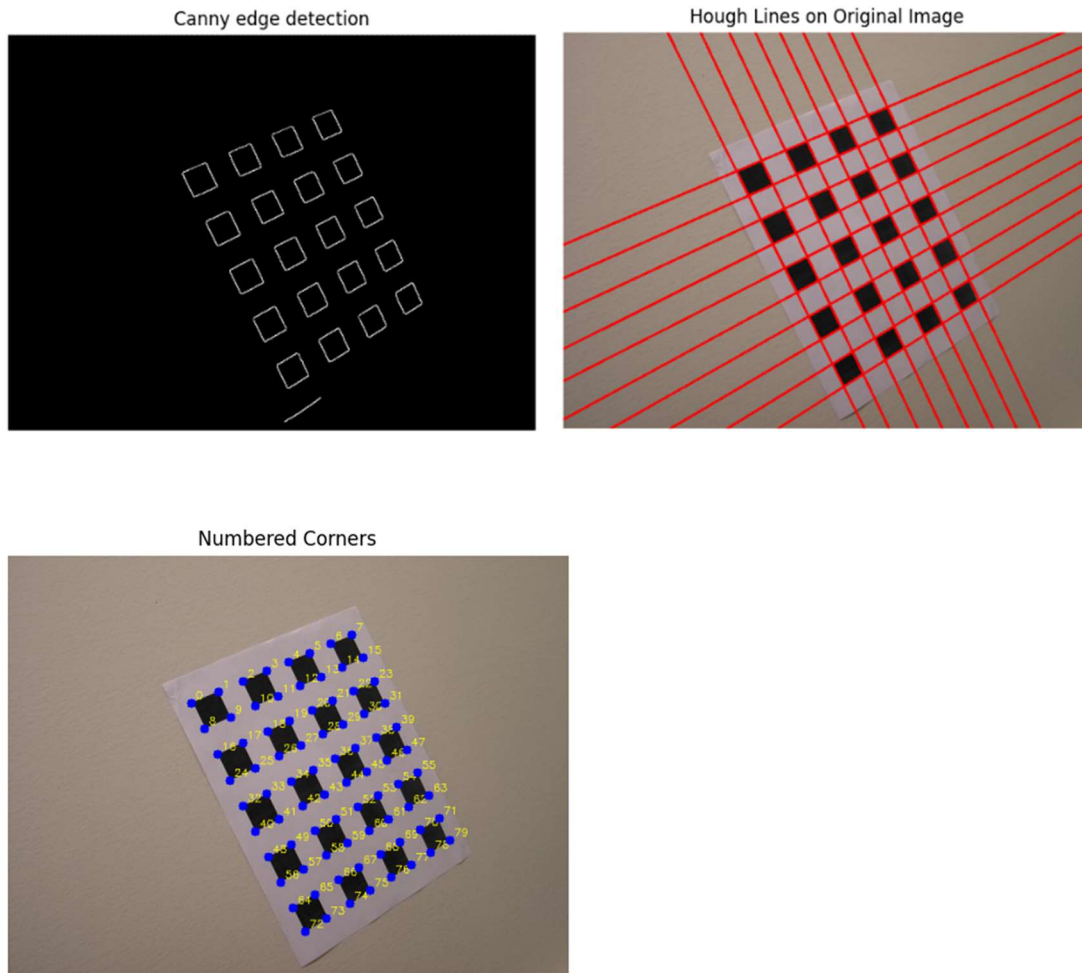


Figure 1: Result from canny edge detection, Hough lines and labeled corners for image 4

The intrinsic matrix K before LM is

$$K = \begin{bmatrix} 0.0202 & -0.0105 & 0.0368 \\ 0 & 0.0415 & 0.0427 \\ 0 & 0 & 1 \end{bmatrix}$$

The rotation and translation parameters of image 4 before LM is

$$[R_4|t_4] = \left[\begin{array}{ccc|c} 0.863 & 0.568 & -0.00001 & 21.91 \\ -0.568 & 0.827 & 0.000015 & 6.609 \\ 0.000016 & -0.000006 & 1 & 0.0016 \end{array} \right]$$




Figure 2: Initial reprojection for image 4

Image 19 Output

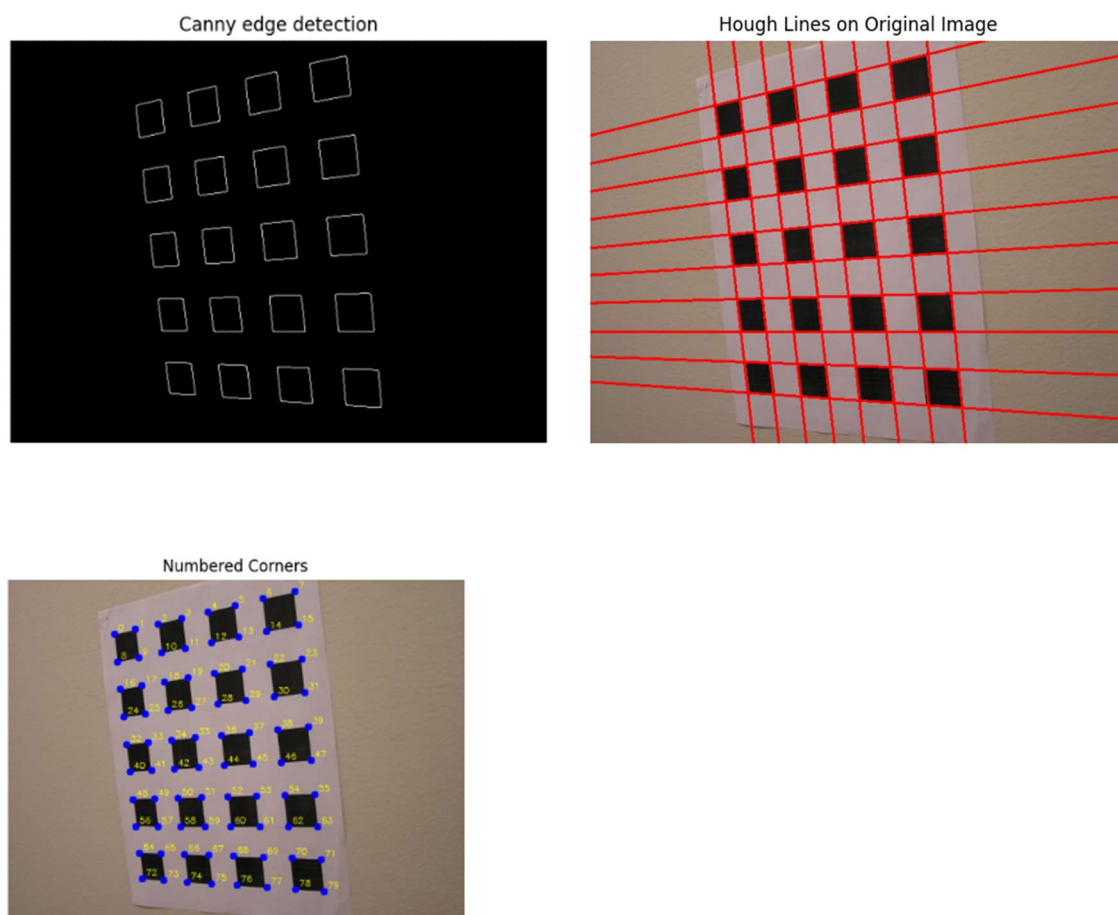


Figure 3: Result from canny edge detection, Hough lines and labeled corners for image 19




Figure 4: Initial reprojection for image 19

The rotation and translation parameters of image 19 before LM is

$$[R_{19}|t_{19}] = \left[\begin{array}{ccc|c} -0.895 & -0.446 & -0.000031 & -27.01 \\ 0.446 & -0.8949 & -0.000025 & -4.53 \\ 0.00004 & -0.000009 & 1 & -0.0028 \end{array} \right]$$

Refinement and radial distortion removal

$$K = \begin{bmatrix} 724.221 & 2.438 & 326.543 \\ 0 & 721.957 & 235.034 \\ 0 & 0 & 1 \end{bmatrix}$$

The rotation and translation parameters of image 4 after LM is

$$[R_4|t_4] = \left[\begin{array}{ccc|c} 0.826 & 0.430 & -0.364 & -8.44 \\ -0.416 & 0.901 & 0.123 & -4.89 \\ 0.381 & 0.05 & 0.923 & 52.89 \end{array} \right]$$




Figure 5: Reprojection after LM and Reprojection after radial distortion removal for image 4

The rotation and translation parameters of image 19 after LM is

$$[R_{19}|t_{19}] = \begin{bmatrix} -0.876 & -0.104 & 0.471 & 10.94 \\ 0.0724 & -0.994 & -0.085 & 9.75 \\ 0.477 & -0.04 & 0.878 & -44.38 \end{bmatrix}$$




Figure 6: Reprojection after LM and Reprojection after radial distortion removal for image 19

Radial distortion parameters

$$k = [-1.06 \times 10^{-7}, \quad -2.16 \times 10^{-12}]$$

Table 1. Comparison of mean and variance before LM, after LM and with radial distortion correction (custom dataset)

	Before LM	After LM	Radial distortion
Error mean (Image 4)	94.816	0.739	0.7301
Error variance (Image 4)	3076.916	0.138	0.16
Error mean (Image 19)	76.326	0.981	0.861
Error variance (Image 19)	1117.62	0.275	0.271

Camera pose

3D plot showing camera poses




Figure 7: 3D view of all camera pose

My dataset

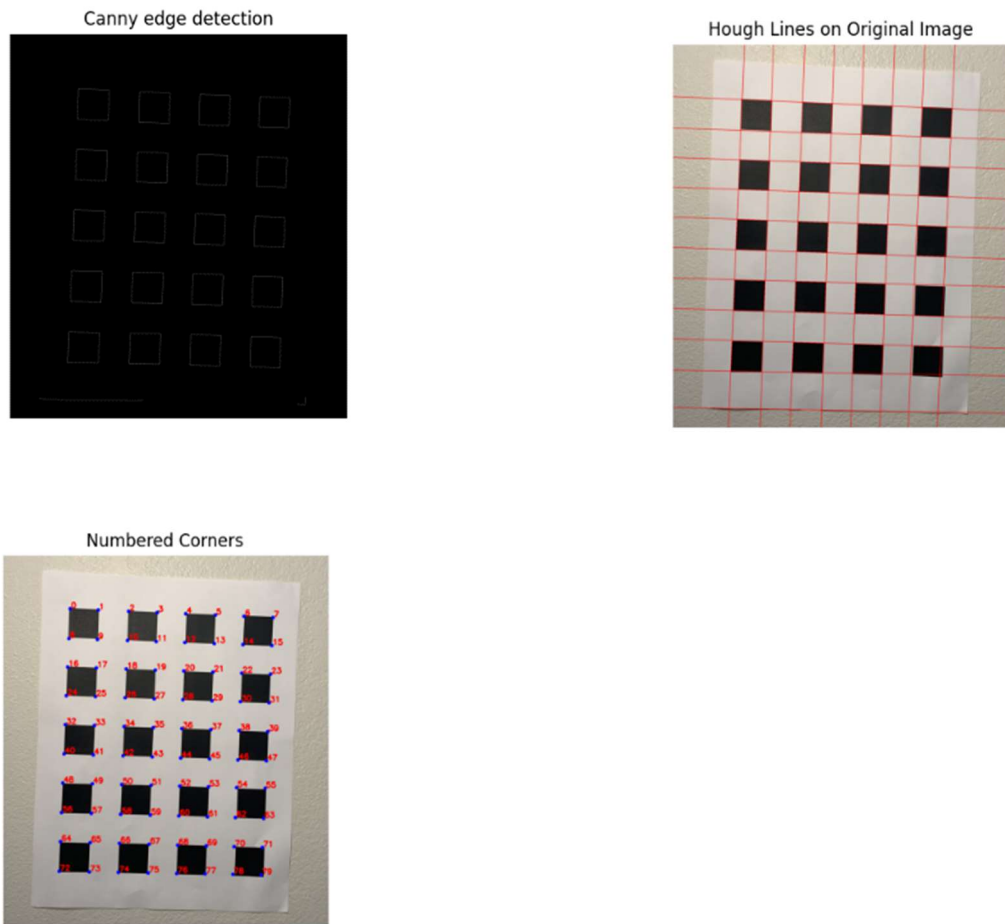


Figure 8: Result from canny edge detection, Hough lines and labeled corners for image 4

The intrinsic matrix K before LM is

$$K = \begin{bmatrix} 0.065 & -0.015 & 0.033 \\ 0 & 0.068 & 0.038 \\ 0 & 0 & 1 \end{bmatrix}$$

The rotation and translation parameters of image 4 before LM is

$$[R_4|t_4] = \begin{bmatrix} -0.997 & -0.071 & -0.000001 & -6.34 \\ 0.0707 & -0.997 & 0.0000003 & -4.12 \\ -0.000001 & -0.000002 & 1 & -0.001 \end{bmatrix}$$






Figure 9: Initial reprojection for image 4

Canny edge detection



Hough Lines on Original Image



Numbered Corners




Figure 10: Result from canny edge detection, Hough lines and labeled corners for image 19

The rotation and translation parameters of image 19 before LM is

$$[R_{19}|t_{19}] = \left[\begin{array}{ccc|c} -0.992 & 1.223 & 0.000003 & 4.83 \\ -1.223 & 0.992 & 0.0000007 & 3.97 \\ -0.000003 & -0.0000004 & 1 & 0.001 \end{array} \right]$$




Figure 11: Initial reprojection for image 4

Refinement and radial distortion removal

$$K = \begin{bmatrix} 341.08 & -2.978 & 993.28 \\ 0 & 338.74 & 569.31 \\ 0 & 0 & 1 \end{bmatrix}$$

The rotation and translation parameters of image 4 after LM is

$$[R_4|t_4] = \left[\begin{array}{ccc|c} -0.998 & 0.038 & -0.531 & 9.25 \\ -0.039 & -1. & 0.011 & 4.022 \\ -0.052 & 0.013 & 0.999 & -50.57 \end{array} \right]$$





Figure 12: Reprojection after LM and Reprojection after radial distortion removal for image 4

The rotation and translation parameters of image 19 after LM is

$$[R_{19}|t_{19}] = \left[\begin{array}{ccc|c} 0.99 & -0.0035 & 0.141 & -12.44 \\ 0.00094 & 0.998 & -0.018 & -5.36 \\ -0.14 & -0.018 & 1 & 55.18 \end{array} \right]$$





Figure 13: Reprojection after LM and Reprojection after radial distortion removal for image 19

Radial distortion parameters

$$k = [4.34 \times 10^{-9}, \quad -6.58 \times 10^{-16}]$$

Table 2. Comparison of mean and variance before LM, after LM and with radial distortion correction (my dataset)

	Before LM	After LM	Radial distortion
Error mean (Image 4)	105.14	2.93	2.97
Error variance (Image 4)	2236.8	3.38	2.17
Error mean (Image 19)	92.19	3.45	3.35
Error variance (Image 19)	1531.57	3.099	2.62

Fixed image vs ground truth




Figure 14: Fixed image vs ground truth

Camera pose

3D plot showing camera poses




Figure 15: 3D view of all camera pose

Code

```
Run Cell | Run Below | Debug Cell
1  # %%
2  import cv2
3  import os
4  import matplotlib.pyplot as plt
5  import numpy as np
6  from scipy.optimize import least_squares
7  import matplotlib.patches as patches
8  import re
9  from mpl_toolkits.mplot3d.art3d import Poly3DCollection
10 from matplotlib import cm
11
Run Cell | Run Above | Debug Cell
12 # %%
13 def group_lines_by_rho(lines, threshold=10):
14     grouped_lines = []
15     lines = sorted(lines, key=lambda x: abs(x[0])) # Sort by |rho|
16
17     current_group = [lines[0]]
18     for line in lines[1:]:
19         if abs(line[0] - current_group[-1][0]) < threshold:
20             current_group.append(line)
21         else:
22             grouped_lines.append(current_group)
23             current_group = [line]
24
25     grouped_lines.append(current_group)
26     return grouped_lines
27
28 def average_line(group):
29     avg_rho = np.mean([line[0] for line in group])
30     avg_theta = np.mean([line[1] for line in group])
31     return avg_rho, avg_theta
32
33 def find_intersection(line1, line2):
34     rho1, theta1 = line1
35     rho2, theta2 = line2
36     A = np.array([
37         [np.cos(theta1), np.sin(theta1)],
38         [np.cos(theta2), np.sin(theta2)]
39     ])
40     b = np.array([rho1, rho2]) # Make this a 1D array instead of a 2D array
41     solution = np.linalg.solve(A, b) # This will be a 1D array, not a 2D array
42     x, y = solution[0], solution[1]
43     return int(round(x)), int(round(y))
44
45 def order_coordinates(corners):
46     # corners = sorted(corners, key=lambda point: point[0])
```



```

47     corners=sorted(corners, key=lambda point: (point[0], point[1]))
48
49     # Group coordinates by their x-value into columns
50     tolerance = 10 # Set a tolerance for grouping points within the same "column"
51     columns = []
52     current_column = [corners[0]]
53
54     for i in range(1, len(corners)):
55         if abs(corners[i][0] - current_column[-1][0]) < tolerance:
56             # If x-value is close enough, it belongs to the same column
57             current_column.append(corners[i])
58         else:
59             # New column
60             columns.append(current_column)
61             current_column = [corners[i]]
62
63     # Add the last column
64     if current_column:
65         columns.append(current_column)
66
67     # Sort each column by y-value to get top-to-bottom order within each column
68     for col in columns:
69         col.sort(key=lambda coord: coord[1])
70
71     # Flatten the columns list to get the final left-to-right, top-to-bottom order
72     corners = [coord for col in columns for coord in col]
73     corners=[[x, y] for x, y in corners]
74     return corners
75
76 def remove_close_points(points, threshold):
77
78     new_points = []
79     prev_point = None
80
81     for point in points:
82         if prev_point is None or np.linalg.norm(np.array(point) - np.array(prev_point)) > threshold:
83             new_points.append(point)
84             prev_point = point
85
86     return new_points
87
88
89 def generate_world_coord(box_size=2.4, rows=5, cols=4):
90
91     # To store the coordinates of all black box corners

```

```

92     box_corners = []
93
94     # Generate coordinates with upper-left origin
95     for row in range(rows):
96         for col in range(cols):
97             # Calculate the starting position of each black box (upper-left corner)
98             x_start = col * 2 * box_size # Every other column
99             y_start = row * 2 * box_size # Positive for each row downwards from the upper-left origin
100
101             # Each black box has four corners, rounded to 1 decimal place
102             corners = [
103                 (round(x_start, 1), round(y_start, 1)), # Upper-left
104                 (round(x_start + box_size, 1), round(y_start, 1)), # Upper-right
105                 (round(x_start + box_size, 1), round(y_start + box_size, 1)), # Bottom-right
106                 (round(x_start, 1), round(y_start + box_size, 1)), # Bottom-left
107             ]
108
109             box_corners.extend(corners)
110     box_corners = [[x, y] for x, y in box_corners]
111     return order_coordinates(box_corners)
112 world_coord = generate_world_coord()
113
114
115 # Function to solve for the last row of V in SVD
116 def linear_least_squares(A):
117     _, _, vh = np.linalg.svd(A, full_matrices=False)
118     return vh[-1]
119
120 # Function to compute homography matrix from domain and range points
121 def calculate_homography(domain_points, range_points):
122     A = []
123
124     # Build matrix A with corresponding points from domain and range
125     for (x, y), (x_prime, y_prime) in zip(domain_points, range_points):
126         A.append([0, 0, 0, -x, -y, -1, y_prime * x, y_prime * y, y_prime])
127         A.append([x, y, 1, 0, 0, 0, -x_prime * x, -x_prime * y, -x_prime])
128
129     # Perform linear least squares to solve for homography matrix
130     H = linear_least_squares(np.array(A))
131
132     # Reshape the result to a 3x3 matrix
133     return H.reshape((3, 3))
134
135 def point_pre_process(points, threshold=10, max_points=80):
136     new_points = []
137     prev_point = None

```

```

138
139     for point in points:
140         if prev_point is None or np.linalg.norm(np.array(point) - np.array(prev_point)) > threshold:
141             new_points.append(point)
142             prev_point = point
143
144     if len(new_points) > max_points:
145         return new_points[:max_points]
146
147     return new_points
148
149
150 def group_lines_by_rho_and_theta(lines, rho_threshold=10, theta_threshold=np.pi / 36):
151     """
152     Groups lines based on proximity in both rho and theta.
153     `rho_threshold` defines the maximum allowed difference in rho.
154     `theta_threshold` defines the maximum allowed difference in theta.
155     """
156     grouped_lines = []
157     lines = sorted(lines, key=lambda x: (x[0], x[1])) # Sort by rho and theta
158
159     current_group = [lines[0]]
160     for line in lines[1:]:
161         if (abs(line[0] - current_group[-1][0]) < rho_threshold and
162             abs(line[1] - current_group[-1][1]) < theta_threshold):
163             current_group.append(line)
164         else:
165             grouped_lines.append(current_group)
166             current_group = [line]
167
168     grouped_lines.append(current_group)
169     return grouped_lines
170
171 def average_line(group):
172     """
173     Averages rho and theta for a group of lines.
174     """
175     avg_rho = np.mean([line[0] for line in group])
176     avg_theta = np.mean([line[1] for line in group])
177     return (avg_rho, avg_theta)
178
179
180
181
182 def get_homography_intersection(dir, threshold1=50, threshold2=150, hough_threshold=50):
183     all_homographies=[]
184     all_intersections=[]

```

```

185 all_edges=[]
186 all_corner_images=[]
187 all_hough_images=[]
188 images=[dir+filename for filename in os.listdir(dir) if filename.endswith(".jpg")]
189 images=sorted(images, key=lambda x: int(re.search(r'Pic_(\d+)', x).group(1)))
190
191
192
193 for img in images:
194     image = cv2.imread(img) # Replace 'checkerboard.png' with your image path
195     hough_image=image.copy()
196     corner_image=image.copy()
197
198     # Convert to grayscale
199     gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
200
201     # adjusted_image = cv2.convertScaleAbs(gray, alpha=1.5, beta=50)
202     # Apply Gaussian blur
203     # blurred = cv2.equalizeHist(gray)
204     # blurred = cv2.GaussianBlur(blurred, (5, 5), 0)
205
206     blurred = cv2.GaussianBlur(gray, (5, 5), 0)
207
208
209
210     # Apply Canny edge detection
211     edges = cv2.Canny(blurred, threshold1=threshold1, threshold2=threshold2) # Adjust thresholds as needed
212
213
214     lines = cv2.HoughLines(edges, rho=1, theta=np.pi/180, threshold=hough_threshold)
215
216     vertical_lines = []
217     horizontal_lines = []
218
219     for line in lines:
220         rho, theta = line[0]
221         if abs(theta) < np.pi / 4 or abs(theta - np.pi) < np.pi / 4:
222             vertical_lines.append((rho, theta))
223         elif abs(theta - np.pi / 2) < np.pi / 4:
224             horizontal_lines.append((rho, theta))
225
226     grouped_vertical_lines = group_lines_by_rho_and_theta(vertical_lines, rho_threshold=70, theta_threshold=np.pi/36)
227     grouped_horizontal_lines = group_lines_by_rho_and_theta(horizontal_lines, rho_threshold=70, theta_threshold=np.pi/36)
228
229     average_vertical_lines = [average_line(group) for group in grouped_vertical_lines]
230     average_horizontal_lines = [average_line(group) for group in grouped_horizontal_lines]
231

```

```

232     for rho, theta in average_vertical_lines + average_horizontal_lines:
233         a = np.cos(theta)
234         b = np.sin(theta)
235         x0 = a * rho
236         y0 = b * rho
237         x1 = int(x0 + 3000 * (-b))
238         y1 = int(y0 + 3000 * (a))
239         x2 = int(x0 - 3000 * (-b))
240         y2 = int(y0 - 3000 * (a))
241         cv2.line(hough_image, (x1, y1), (x2, y2), (0, 0, 255), 2)
242
243     corners = []
244     for h_line in average_horizontal_lines:
245         for v_line in average_vertical_lines:
246             corner = find_intersection(h_line, v_line)
247             corners.append(corner)
248
249     corners = point_pre_process(corners)
250
251     homography = calculate_homography(world_coord, corners)
252
253     all_intersections.append(corners)
254     all_homographies.append(homography)
255     all_hough_images.append(hough_image)
256     all_corner_images.append(corner_image)
257     all_edges.append(edges)
258
259     return all_homographies, all_intersections, all_edges, all_corner_images, all_hough_images
260
261
262
263
264
265
266 def plot_edges_hough_corner(corners, edges, corner_image, hough_image, img_no):
267     # corners=intersection,
268
269     for idx, point in enumerate(corners):
270         # print(point)
271         cv2.circle(corner_image, point, 10, (255, 0, 0), -1) # Draw the point
272         cv2.putText(corner_image, str(idx), (point[0] + 5, point[1] - 5),
273                    cv2.FONT_HERSHEY_SIMPLEX, 1.4, (0, 0, 255), 5) # Label with numbers starting from 0
274
275     # # Display the results
276     plt.imshow(edges, cmap='gray'), plt.title('Edges')
277     plt.title('Canny edge detection')
278
279     plt.axis('off')
280     plt.savefig(f'{img_no} Canny edge detection')
281     # plt.show()
282     # plt.subplot(122), plt.imshow(cv2.cvtColor(original_image, cv2.COLOR_BGR2RGB)), plt.title('Hough Line Segments')
283
284     plt.imshow(cv2.cvtColor(hough_image, cv2.COLOR_BGR2RGB))
285     plt.title('Hough Lines on Original Image')
286     plt.axis('off')
287     plt.savefig(f'{img_no} Hough Lines on Original Image')
288     # plt.show()
289
290     plt.imshow(cv2.cvtColor(corner_image, cv2.COLOR_BGR2RGB))
291     plt.title('Numbered Corners')
292     plt.axis('off')
293     plt.savefig(f'{img_no} Numbered Corners')
294     # plt.show()
295
296     all_homographies, all_intersections, edge_img, corners_img, hough_img = get_homography_intersection('HmB-Files/Pattern/', threshold1=200, threshold2=500, hough_threshold=100)
297     plot_edges_hough_corner(all_intersections[3], edge_img[3], corners_img[3], hough_img[3], 4)
298     plot_edges_hough_corner(all_intersections[18], edge_img[18], corners_img[18], hough_img[18], 19)
299
300
301 Run Cell | Run Above | Debug Cell
302 # %%
303 # Image of the absolute conic
304 def find_omega(Hs):
305     V = []
306     for H in Hs:
307         h11, h12, h13 = H[0, 0], H[0, 1], H[0, 2]
308         h21, h22, h23 = H[1, 0], H[1, 1], H[1, 2]
309
310         # Append two constraints derived from each homography
311         V.append([h11 * h21, h11 * h22 + h12 * h21, h12 * h22, h13 * h21 + h11 * h23, h13 * h22 + h12 * h23, h13 * h23])
312         V.append([h11 ** 2 - h21 ** 2, 2 * (h11 * h12 - h21 * h22), h12 ** 2 - h22 ** 2, 2 * (h13 * h11 - h23 * h21), 2 * (h13 * h12 - h23 * h22), h13 ** 2 - h23 ** 2])
313
314     # Solve for omega components
315     b = linear_least_squares(np.array(V))
316
317     # Create the symmetric matrix omega from b vector
318     omega = np.array([
319         [b[0], b[1], b[3]],
320         [b[1], b[2], b[4]],
321         [b[3], b[4], b[5]]
322     ])
323
324     return omega
325
326 absolute_conic = find_omega(all_homographies)

```

```

Run Cell | Run Above | Debug Cell
326 # %%
327 def find_intrinsic(omega):
328     # Extract terms from omega matrix
329     x0 = (omega[0, 1] * omega[0, 2] - omega[0, 0] * omega[1, 2]) / (omega[0, 0] * omega[1, 1] - omega[0, 1] ** 2)
330     # Calculate lambda as a scaling factor
331     lambda_ = omega[2, 2] - ((omega[0, 2]**2 + x0 * (omega[0, 1] * omega[0, 2] - omega[0, 0] * omega[1, 2])) / omega[0, 0])
332     # Calculate alpha values as focal lengths
333     alpha_x = np.sqrt(lambda_ / omega[0, 0])
334     alpha_y = np.sqrt(lambda_ * omega[0, 0] / (omega[0, 0] * omega[1, 1] - omega[0, 1] ** 2))
335
336     # Compute skew and y0
337     s = -omega[0, 1] * alpha_x**2 * alpha_y / lambda_
338     y0 = s * x0 / alpha_y - omega[0, 2] * alpha_x**2 / lambda_
339
340     # Construct intrinsic matrix K
341     K = np.array([
342         [alpha_x, s, x0],
343         [0, alpha_y, y0],
344         [0, 0, 1]
345     ])
346
347     return K
348 intrinsic_params=find_intrinsic(absolute_conic)
349
350 print('Intrinsic param before LM',intrinsic_params)
351
352
353
Run Cell | Run Above | Debug Cell
354 # %%
355 def find_extrinsic(Hs, K):
356
357     Rs = []
358     ts = []
359
360     K_inv = np.linalg.inv(K)
361
362     for H in Hs:
363         # Extract the first two columns of the homography matrix and compute r1 and r2
364         r1 = np.dot(K_inv, H[:, 0])
365         r2 = np.dot(K_inv, H[:, 1])
366
367         # Scaling factor to normalize r1
368         scaling = 1.0 / np.linalg.norm(r1)
369         r1 *= scaling
370         r2 *= scaling
371

```

```

372     # Calculate the translation vector
373     t = scaling * np.dot(K_inv, H[:, 2])
374
375     # Compute r3 as the cross product of r1 and r2 to ensure orthogonality
376     r3 = np.cross(r1, r2)
377
378     # Stack r1, r2, and r3 to form the initial rotation matrix
379     R = np.column_stack((r1, r2, r3))
380
381     # Orthogonalize R using SVD to enforce a valid rotation matrix
382     u, _, vh = np.linalg.svd(R)
383     conditioned_R = np.dot(u, vh)
384
385     # Append the conditioned rotation matrix and translation vector to the lists
386     Rs.append(conditioned_R)
387     ts.append(t)
388
389     return Rs, ts
390
391 rotation_params, translation_params = find_extrinsic(all_homographies, intrinsic_params)
392 print('-----')
393 print('Rotation 4 before LM', rotation_params[3], 'translation', translation_params[3])
394
395 print('Rotation 19 before LM', rotation_params[18], 'translation', translation_params[18])
396 # print('-----')
397
398 Run Cell | Run Above | Debug Cell
399 # %%
400 # Extract camera parameters including rotation and translation
401 def camera_parameters(K, Rs, ts, radio_distortion=False):
402     # Initialize parameter array with intrinsic parameters from K
403     params = [K[0, 0], K[0, 1], K[0, 2], K[1, 1], K[1, 2]]
404
405     for R, t in zip(Rs, ts):
406         # Calculate rotation angle
407         phi = np.arccos((np.trace(R) - 1) / 2)
408
409         # Calculate Rodrigues vector w (rotation vector)
410         if np.sin(phi) != 0:
411             w = phi / (2 * np.sin(phi)) * np.array([R[2, 1] - R[1, 2], R[0, 2] - R[2, 0], R[1, 0] - R[0, 1]])
412         else:
413             w = np.array([0, 0, 0]) # Edge case when phi = 0 (no rotation)
414
415         # Append rotation and translation to parameters
416         params.extend(w)
417         params.extend(t)

```



```

418     # Add radial distortion parameters if required
419     if radio_distortion:
420         params.extend([0, 0])
421
422     return np.array(params)
423
424 params=camera_parameters(intrinsic_params, rotation_params, translation_params)
425
426 Run Cell | Run Above | Debug Cell
427 # %%
428
429 def apply_homography(H, domain_points):
430     # Add homogeneous coordinate (1) to each point
431     domain_points=np.array(domain_points)
432     num_points = domain_points.shape[0]
433     homo_domain_points = np.hstack((domain_points, np.ones((num_points, 1))))
434
435     # Apply homography transformation
436     range_points = (H @ homo_domain_points.T).T
437
438     # Normalize by the third (homogeneous) coordinate
439     range_points /= range_points[:, 2, np.newaxis]
440     return range_points[:, :2] # Return only x and y coordinates
441
442 def reconstruct_R(params):
443     # Intrinsic matrix K
444     K = np.array([[params[0], params[1], params[2]],
445                  [0, params[3], params[4]],
446                  [0, 0, 1]])
447
448     Rs = []
449     ts = []
450
451     # Iterating through the parameters to calculate rotation and translation
452     for i in range(5, 6 * ((len(params) - 5) // 6), 6):
453         w = np.array(params[i:i+3])
454         t = np.array(params[i+3:i+6])
455
456         # Calculating rotation matrix using Rodrigues' formula
457         phi = np.linalg.norm(w)
458         if phi == 0:
459             R = np.eye(3) # No rotation if phi is zero
460         else:
461             w_matrix = np.array([[0, -w[2], w[1]],
462                                [w[2], 0, -w[0]],
463                                [-w[1], w[0], 0]])
464             R = (nn_evp(1) +

```



```

464         (np.sin(phi) / phi) * w_matrix +
465         ((1 - np.cos(phi)) / (phi ** 2)) * np.dot(w_matrix, w_matrix))
466
467     # Append rotation and translation matrices to the lists
468     Rs.append(R)
469     ts.append(t)
470
471     return K, Rs, ts
472
473
474
475 # Remove radial distortion from projected points
476 def remove_radial_distortion(projected_points, k1, k2, x0, y0):
477     x, y = projected_points[:, 0], projected_points[:, 1]
478     # Calculate radial distance squared
479     r2 = (x - x0) ** 2 + (y - y0) ** 2
480     # Apply distortion correction
481     x_rad = x + (x - x0) * (k1 * r2 + k2 * r2 ** 2)
482     y_rad = y + (y - y0) * (k1 * r2 + k2 * r2 ** 2)
483     return np.column_stack((x_rad, y_rad))
484
485 def cost_function(params, all_intersec_points, world_coord, radio_distortion=False):
486     if radio_distortion:
487         # print(params[:-10].shape)
488         K, Rs, ts = reconstruct_R(params[:-2])
489         k1, k2 = params[-2], params[-1]
490         x0, y0 = params[2], params[4]
491     else:
492         K, Rs, ts = reconstruct_R(params)
493
494     all_projected_points = []
495
496     for R, t in zip(Rs, ts):
497         H = np.dot(K, np.column_stack((R[:, 0], R[:, 1], t)))
498         projected_points = apply_homography(H, world_coord)
499
500         if radio_distortion:
501             projected_points = remove_radial_distortion(projected_points, k1, k2, x0, y0)
502
503         all_projected_points.append(projected_points)
504
505     # print(all_projected_points)
506     all_projected_points = np.concatenate(all_projected_points, axis=0)
507     all_intersec_points = np.concatenate(all_intersec_points, axis=0)
508     diff = all_intersec_points - all_projected_points
509

```

```

501         projected_points = remove_radio_distortion(projected_points, k1, k2, x0, y0)
502
503
504         all_projected_points.append(projected_points)
505
506         # print(all_projected_points)
507         all_projected_points = np.concatenate(all_projected_points, axis=0)
508         all_intersec_points = np.concatenate(all_intersec_points, axis=0)
509         diff = all_intersec_points - all_projected_points
510         return diff.flatten()
511
512
513 # Run least squares optimization
514 res_lsq = least_squares(
515     cost_function,
516     params,
517     method='lm', # Levenberg-Marquardt method
518     args=[all_intersections, world_coord]
519 )
520
521 Run Cell | Run Above | Debug Cell
522 # %%
523 refined_K, refined_R, refined_t=reconstruct_R(res_lsq.x)
524
525 print('After refinement camera matrix',refined_K,'Rotation 4 before LM',refined_R[3], 'translation', refined_t[3])
526 print('After refinement Rotation 19 before LM',refined_R[18], 'translation', refined_t[18])
527 # print('-----')
528
529 Run Cell | Run Above | Debug Cell
530 # %%
531 def error_per_image(diff):
532     dx = diff[:, 0]
533     dy = diff[:, 1]
534     distance = np.sqrt(dx**2 + dy**2)
535     mean = np.mean(distance)
536     var = np.var(distance)
537     return mean, var
538
539 def reprojection(world_coord, params, all_intersec_points, rad_distortion=False, img_idx=0, save_img=False, output_name=None):
540     if rad_distortion:
541         K, Rs, ts = reconstruct_R(params[:-2])
542         k1 = params[-2]
543         k2 = params[-1]
544         x0 = params[2]
545         y0 = params[4]
546     else:
547
548         all_projected_points = []
549         K, Rs, ts = reconstruct_R(params)
550         for R, t in zip(Rs, ts):
551             H = np.matmul(K, np.column_stack((R[:, 0], R[:, 1], t)))
552             projected_points = apply_homography(H, world_coord)
553
554             if rad_distortion:
555                 projected_points = remove_radio_distortion(projected_points, k1, k2, x0, y0)
556
557             all_projected_points.append(projected_points)
558
559         if img_idx >= 0:
560             diff = all_intersec_points[img_idx] - all_projected_points[img_idx]
561             mean, var = error_per_image(diff)
562
563         if save_img:
564             img_path = 'HW8-Files/Pattern/Pic_' + str(img_idx+1) + '.jpg'
565             img = plt.imread(img_path)
566
567             # Set up the plot with the image
568             fig, ax = plt.subplots()
569             ax.imshow(img)
570
571             # Draw circles at each projected point
572             color='lime'
573             if output_name == 'radio_distortion':
574                 color='red'
575
576             # Plot original target view points in blue
577             for idx, point in enumerate(all_intersec_points[img_idx]):
578                 circle = patches.Circle((point[0], point[1]), radius=3, color='blue', fill=True)
579
580                 ax.add_patch(circle)
581                 ax.text(point[0] + 5, point[1] - 5, str(idx), fontsize=8, color='yellow',
582                     ha='left', va='top')
583
584             for point in all_projected_points[img_idx]:
585                 circle = patches.Circle((point[0], point[1]), radius=3, color=color, fill=True)
586                 ax.add_patch(circle)
587
588             # Remove axis ticks for a cleaner image
589             ax.axis('off')
590
591             # Save the image
592             output_path = 'Pic_' + str(img_idx+1) + '_' + output_name + '_reproject.jpg'

```

```

592         output_path = 'Pic_' + str(img_idx+1) + '_' + output_name + '_reproject.jpg'
593         plt.savefig(output_path, bbox_inches='tight', pad_inches=0)
594         plt.close(fig)
595
596     return mean, var
597
Run Cell | Run Above | Debug Cell
598 # %%
599 mean, var=reprojection(world_coord, params, all_intersections, img_idx=3, save_img=True, output_name='init')
600 print('Img 4, init mean and variance',mean, var)
601
602 mean, var=reprojection(world_coord, res_lsq.x, all_intersections, img_idx=3, save_img=True, output_name='refined')
603 print('Img 4, refined mean and variance',mean, var)
604
605 mean, var=reprojection(world_coord, params, all_intersections, img_idx=18, save_img=True, output_name='init')
606 print('Img 19, init mean and variance',mean, var)
607
608 mean, var=reprojection(world_coord, res_lsq.x, all_intersections, img_idx=18, save_img=True, output_name='refined')
609 print('Img 19, refined mean and variance',mean, var)
610 print('-----')
611
612
Run Cell | Run Above | Debug Cell
613 # %%
614 params_radio_distortion=camera_parameters(intrinsic_params, rotation_params, translation_params, radio_distortion=True)
615
616
617 # Run least squares optimization
618 res_lsq_radio_distortion = least_squares(
619     cost_function,
620     params_radio_distortion,
621     method='lm', # Levenberg-Marquardt method
622     args=[all_intersections, world_coord, True]
623 )
624
625 print('Radial distortion is',res_lsq_radio_distortion.x[-2:])
626
Run Cell | Run Above | Debug Cell
627 # %%
628 mean, var=reprojection(world_coord, res_lsq_radio_distortion.x, all_intersections, rad_distortion=True, img_idx=3, save_img=True, output_name='radio_distortion')
629 print('Img 4, radio distortion mean and variance',mean, var)
630
631 mean, var=reprojection(world_coord, res_lsq_radio_distortion.x, all_intersections, rad_distortion=True, img_idx=18, save_img=True, output_name='radio_distortion')
632 print('Img 19, radio distortion mean and variance',mean, var)
633 print('-----')
634
635
636 # Function to plot camera pose
637 def plot_camera(ax, R, t, scale=5, color='gray',plane_color='cyan'):
638     # Define camera frame axes in the camera's local coordinates
639     X_cam = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]).T * scale
640
641     # Rotate and translate camera frame axes to world coordinates
642     X_world = R @ X_cam + t.reshape(-1, 1)
643
644     # Plot the camera center
645     ax.scatter(t[0], t[1], t[2], color=color, s=50)
646
647     # Plot the axes
648     ax.quiver(t[0], t[1], t[2], X_world[0, 0]-t[0], X_world[1, 0]-t[1], X_world[2, 0]-t[2], color='r', arrow_length_ratio=0.1)
649     ax.quiver(t[0], t[1], t[2], X_world[0, 1]-t[0], X_world[1, 1]-t[1], X_world[2, 1]-t[2], color='g', arrow_length_ratio=0.1)
650     ax.quiver(t[0], t[1], t[2], X_world[0, 2]-t[0], X_world[1, 2]-t[1], X_world[2, 2]-t[2], color='b', arrow_length_ratio=0.1)
651
652     # Draw the camera plane
653     plane_points = np.array([[0.5, 0.5, 1], [0.5, -0.5, 1], [-0.5, -0.5, 1], [-0.5, 0.5, 1]]).T * scale
654     plane_world = R @ plane_points + t.reshape(-1, 1)
655
656     verts = [list(zip(plane_world[0, :], plane_world[1, :], plane_world[2, :]))]
657     ax.add_collection3d(Poly3DCollection(verts, color=plane_color, alpha=0.2))
658
659 # Set up the figure and 3D axis
660 fig = plt.figure(figsize=(8, 6))
661 ax = fig.add_subplot(111, projection='3d')
662
663 K, Rs, ts = reconstruct_R(res_lsq.x)
664 # print(Rs[0])
665 # print('ts',ts[0])
666
667 # print('Rs',Rs[3])
668 # print('ts',ts[3])
669
670 camera_poses=[(Rs[0], ts[0]), (Rs[3], ts[3])]
671 # Plot each camera
672 cam=[]
673 for r,t in zip(Rs, ts):
674     new=(r,t)
675     cam.append(new)
676
677 colormap = cm.get_cmap('viridis', len(cam))
678
679
680 for i, (R, t) in enumerate(cam):
681     color = colormap(i)
682     plot_camera(ax, R, t, plane_color=color)
683

```

```
684 # Set plot limits and labels
685 ax.set_xlim(-20, 20)
686 ax.set_ylim(-20, 20)
687 ax.set_zlim(-60, 60)
688 ax.set_xlabel("X")
689 ax.set_ylabel("Y")
690 ax.set_zlabel("Z")
691 ax.set_title("3D plot showing camera poses")
692
693 plt.show()
694
```